

✓ XGBoost for Predicting Energy Usage in Buildings

Notebook Overview

This notebook demonstrates:


1. **Data Preprocessing:**
 - Cleaning, aligning, and preparing building-specific data.
2. **Model Evaluation:**
 - Using a trained XGBoost model to predict energy consumption.
3. **Visualization:**
 - Comparing actual and predicted values with interactive plots.

Use cases:

- **Model Validation:** Ensure predictions align with ground truth data.
- **Error Analysis:** Identify where and why predictions deviate from reality.
- **Presentation:** Generate clear visuals for stakeholder communication.

```
1 import pandas as pd
2 import numpy as np
3 import os
4 import matplotlib.pyplot as plt
5 import re
6 import json
7 import gc
8 import psutil
9 import xgboost as xgb
10 import random
11 import joblib
12 import plotly.express as px
13
14 from multiprocessing import Pool, cpu_count
15 from concurrent.futures import ProcessPoolExecutor, as_completed
16 from functools import partial
17
18 from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
19 from sklearn.model_selection import RandomizedSearchCV
```

```
1 # if you are using Google Colab, you can easily mount the drive and access
2 # the data.
3 from google.colab import drive
4 drive.mount('/content/drive', force_remount=True)
```

 Mounted at /content/drive

✓ Load the datasets

```
1 # In order for you to use the data, you need to update these paths.
2 PROCESSED = f'/content/drive/MyDrive/Team-Fermata-Energy/processed_data/' # https://drive.google.com/drive/folders,
3 BUILDINGS = f'{PROCESSED}processed_weather_load_w_timestamp/' # https://drive.google.com/drive/folders/1kW3Ip5_xm6l
```

```


1 with open(f'{PROCESSED}subset20_data.json', 'r') as test_train_file:
2     test_train_ids = json.load(test_train_file)
3
4 train_ids = [int(bldg_id.replace('.csv', '')) for bldg_id in test_train_ids['train_bldg_ids']]
5 test_ids = [int(bldg_id.replace('.csv', '')) for bldg_id in test_train_ids['test_bldg_ids']]

```

```

1 df_metadata = pd.read_csv(f'{PROCESSED}md_one_hot_encoded_subset20.csv")
2 df_metadata.head(20)

```



	bldg_id	in.state	in.vintage	in.sqft	in.building_america_climate_zone_Cold	in.building_america_climate_zo
0	105885	10	3	750000.0		0
1	305819	40	2	150000.0		0
2	305934	40	4	350000.0		0
3	317044	40	3	350000.0		0
4	32	1	6	37500.0		0
5	64	1	0	37500.0		0
6	103	1	4	75000.0		0
7	112	1	3	7500.0		0
8	277	1	0	37500.0		0
9	355	1	7	17500.0		0
10	363	1	5	37500.0		0
11	379	1	1	17500.0		0
12	417	1	0	17500.0		0
13	530	1	2	17500.0		0
14	575	1	3	3000.0		0
15	611	1	5	7500.0		0
16	633	1	6	37500.0		0
17	864	1	4	37500.0		0
18	1025	1	4	75000.0		0
19	1215	1	3	3000.0		0

20 rows × 40 columns

✓ If you want to load in the model and just evaluate, you can do so here!

```
1 model = joblib.load(f'{PROCESSED}xgb_model2.pkl')
```

✓ Functions for Training the Model

Training the Model

This function prepares data and trains the XGBoost model for energy consumption prediction.

Steps:

1. Train the model: Set hyperparameters, use early stopping, and monitor performance with metrics like RMSE or MAE.
2. Analyze feature importance: Understand key drivers of predictions.

Outputs:

- A trained model

```

1 def preprocess_bldg_optimized(bldg_id, df_metadata):
2     """
3     Preprocesses data for a single building.
4
5     Parameters:
6     - bldg_id (int): Building ID.
7     - df_metadata (DataFrame): Metadata DataFrame.
8
9     Returns:
10    - X (DataFrame): Feature DataFrame.
11    - y (Series): Target variable.
12    """
13    try:
14        # Load CSV with optimized data types
15        df_bldg = pd.read_csv(
16            f"{BUILDINGS}{bldg_id}.csv",
17            dtype={
18                'bldg_id': 'int32',
19                'minute': 'int8',
20                'out_electricity_total_energy_consumption': 'float32',
21                # Add other columns with appropriate types if known
22                # Example:
23                # 'temperature': 'float32',
24                # 'humidity': 'float32',
25                # 'heat_index': 'float32',
26                # 'location': 'category',
27            }
28        )
29
30        # Clean column names
31        df_bldg.columns = [re.sub(r"^[A-Za-z0-9_]+", "_", col) for col in df_bldg.columns]
32
33        # Filter rows where 'minute' == 0
34        df_bldg = df_bldg[df_bldg['minute'] == 0]
35
36        # Merge with metadata
37        bldg_metadata = df_metadata[df_metadata['bldg_id'] == bldg_id]
38        df_bldg = df_bldg.merge(bldg_metadata, on='bldg_id', how='left')
39
40        # Prepare features and target
41        y = df_bldg['out_electricity_total_energy_consumption']
42        X = df_bldg.drop(columns=['out_electricity_total_energy_consumption', 'timestamp', 'bldg_id'])
43
44        return X, y
45    except Exception as e:
46        print(f"Error processing building ID {bldg_id}: {e}")
47        return pd.DataFrame(), pd.Series()
48
49 def train_in_chunks_optimized(df_metadata, train_bldg_ids, best_param):
50     """
51     Trains an XGBoost model on data from multiple buildings using GPU acceleration.
52
53     Parameters:
54     - df_metadata (DataFrame): Metadata DataFrame.
55     - train_bldg_ids (list): List of building IDs for training.
56     - best_param (dict): Best hyperparameters for XGBoost.
57
58     Returns:
59     - model (XGBModel): Trained XGBoost model.
60     - y_preds (Series): Predictions for the training set.
61     """
62     # Prepare training data
63     X_train, y_train = [], []
64     for bldg_id in train_bldg_ids:
65         X, y = preprocess_bldg_optimized(bldg_id, df_metadata)
66         X_train.append(X)
67         y_train.append(y)
68
69     # Concatenate training data
70     X_train = pd.concat(X_train)
71     y_train = pd.concat(y_train)
72
73     # Train XGBoost model
74     model = xgb.XGBModel(
75         params=best_param,
76         data=(X_train, y_train),
77         num_parallel_tree=1,
78         use_cuda_device=-1,
79         verbose=0,
80     )
81     model.fit(X_train, y_train, eval_metric='rmse')
82
83     # Predict on training set
84     y_preds = model.predict(X_train)
85
86     return model, y_preds

```

```

58     returns:
59     - model (XGBRegressor): Trained XGBoost model.
60     """
61     # Determine number of parallel processes
62     n_jobs = max(cpu_count() - 1, 1) # Reserve one core for the system
63     print(f"Using {n_jobs} parallel processes for data preprocessing.")
64
65     # Initialize multiprocessing Pool
66     with Pool(processes=n_jobs) as pool:
67         # Partial function to pass df_metadata
68         func = partial(preprocess_bldg_optimized, df_metadata=df_metadata)
69
70         # Map the preprocessing function to building IDs
71         results = pool.map(func, train_bldg_ids)
72
73     # Filter out any failed preprocessing results
74     results = [res for res in results if not res[0].empty]
75
76     if not results:
77         raise ValueError("No data was successfully preprocessed.")
78
79     # Concatenate all preprocessed data
80     X_all, y_all = zip(*results)
81     X_all = pd.concat(X_all, ignore_index=True)
82     y_all = pd.concat(y_all, ignore_index=True)
83
84     # Clean up intermediate results
85     del results
86     gc.collect()
87
88     print("Starting model training on GPU...")
89
90     # Initialize and train the XGBoost model with GPU support
91     model = xgb.XGBRegressor(
92         tree_method='hist',          # Use 'hist' for faster training with GPU via 'device'
93         device='cuda',              # Specify to use GPU
94         n_jobs=-1,                  # Utilize all available CPU cores for data preprocessing
95         enable_categorical=True,     # Enable categorical feature support
96         **best_param,               # Additional hyperparameters
97         reg_alpha=1.0,
98         reg_lambda=1.0,
99         random_state=42,
100        verbosity=1                  # Set to 1 for basic logging
101    )
102
103    # Fit the model
104    model.fit(X_all, y_all, verbose=True)
105
106    # Monitor memory usage
107    process = psutil.Process()
108    memory_usage = process.memory_info().rss
109    print(f"Memory Usage After Training: {memory_usage / (1024 ** 2):.2f} MB")
110
111    # Clean up to free memory
112    del X_all
113    del y_all
114    gc.collect()
115
116    return model
117

```

```

1 model = xgb.XGBRegressor(
2     tree_method='hist',          # Use 'hist' for faster training with GPU via 'device'
3     device='cuda',              # Specify to use GPU
4     n_jobs=-1,                  # Utilize all available CPU cores for data preprocessing

```

```

5     enable_categorical=True,      # Enable categorical feature support
6     reg_alpha=1.0,
7     reg_lambda=1.0,
8     random_state=42,
9     verbosity=1                  # Set to 1 for basic logging
10 )

```

```

1 best_param = {'subsample': 0.8, 'n_estimators': 300, 'max_depth': 6, 'learning_rate': 0.01, 'colsample_bytree': 0.1}
2 model = train_in_chunks_optimized(df_metadata, train_ids, best_param)

```

➡ Using 11 parallel processes for data preprocessing.
Starting model training on GPU...
Memory Usage After Training: 37627.39 MB

➤ Find best hyperparameters for the model.

I've already run this code so you don't have to.

[] ↳ 2 cells hidden

✓ Save and Test Model

Evaluating the Model

This function tests the trained model on unseen data and calculates performance metrics.

Steps:

1. Load and preprocess test data: Align features with the trained model.
2. Predict and compare: Generate predictions and calculate metrics like SMAPE.
3. Monitor performance: Assess memory usage and model accuracy.

Outputs:

- Evaluation metrics and insights into prediction accuracy.

```

1 # Save the model
2 # joblib.dump(model, PATHGOESHERE)

```

➡ [' /content/drive/MyDrive/Team-Fermata-Energy/processed_data/xgb_model2.pkl']

```

1 df_test_bldg = pd.read_csv(f"{BUILDINGS}32.csv")
2 df_test_bldg.columns = [re.sub(r"^[A-Za-z0-9_]+", "_", col) for col in df_test_bldg.columns]
3 df_test_bldg.columns

```

➡ Index(['timestamp', 'out_electricity_total_energy_consumption', 'Dry_Bulb_Temperature_C', 'Relative_Humidity', 'heat_index', 'minute', 'hour', 'day', 'month', 'is_weekday', 'is_holiday', 'max_load_hourly', 'min_load_hourly', 'max_temp_hourly', 'min_temp_hourly', 'bldg_id'], dtype='object')

```

1 # Pre-create numpy arrays to store the metrics
2 # num_test_ids = len(test_ids) # Replace test_ids with your list of test building IDs
3 smape_values = []
4
5 def calculate_smape(y_true, y_pred):
6     """
7     Calculate Symmetric Mean Absolute Percentage Error (SMAPE).
8

```

```

9     Parameters:
10         y_true: Actual values.
11         y_pred: Predicted values.
12
13     Returns:
14         SMAPE value as a percentage.
15     """
16     denominator = (np.abs(y_true) + np.abs(y_pred)) / 2
17     diff = np.abs(y_true - y_pred)
18     smape = np.mean(diff / denominator) * 100 # Percentage
19     return smape
20
21 def evaluate_model(model, df_metadata, test_ids):
22     """
23     Loops through each test building ID, evaluates the model, and stores metrics.
24
25     Parameters:
26         model: Trained XGBoost model.
27         df_metadata: DataFrame containing building metadata.
28         test_ids: List of test building IDs.
29     """
30     for idx, bldg_id in enumerate(test_ids):
31         # Load and preprocess test building data
32         df_test_bldg = pd.read_csv(f"{BUILDINGS}{bldg_id}.csv")
33         df_test_bldg.columns = [re.sub(r"^[A-Za-z0-9_]+", "_", col) for col in df_test_bldg.columns]
34
35         df_test_bldg = df_test_bldg[df_test_bldg['minute'] == 0]
36
37         # Merge metadata
38         test_bldg_metadata = df_metadata[df_metadata['bldg_id'] == bldg_id]
39         df_test_bldg = df_test_bldg.merge(test_bldg_metadata, on='bldg_id', how='left')
40
41         # Prepare features (X_test) and target (y_test)
42         y_test = df_test_bldg['out_electricity_total_energy_consumption']
43         X_test = df_test_bldg.drop(columns=['out_electricity_total_energy_consumption', 'timestamp', 'bldg_id'])
44
45         # print(y_test.describe())
46         # Predict and evaluate metrics
47         pred = model.predict(X_test)
48         smape = calculate_smape(y_test, pred)
49         smape_values.append(smape)
50
51         # Monitor memory usage
52         process = psutil.Process()
53         memory_usage = process.memory_info().rss
54         print(f"Memory Usage: {memory_usage / (1024 ** 2):.2f} MB")
55
56         # Clean up memory
57         del df_test_bldg, X_test, y_test, pred, process, memory_usage
58         gc.collect()
59
60 # Example usage
61 random.seed(521)
62 evaluate_model(model, df_metadata, random.choices(test_ids, k = 5))
63
64 # performance metrics
65 smape_array = np.array(smape_values)
66
67 print(f'Mean SMAPE: {mean_smape}')

```

```

➦ Memory Usage: 19700.38 MB
Memory Usage: 19700.38 MB
Memory Usage: 19700.38 MB
Memory Usage: 19700.38 MB
Memory Usage: 19700.38 MB

```

Mean SMAPE: 21.48865254136774

✓ Visualizations and Misc

Align Features for Model Compatibility

The `align_features` function ensures that the test data matches the trained model's expected input format. This is critical for avoiding feature name mismatches, which occur when:

- The test data contains columns that were not part of the training data.
- The test data is missing columns present during model training.

Key steps:

1. **Add Missing Columns:** Columns that are in the model's feature list but not in the test data are added with default values (e.g., 0 for one-hot encoded features).
2. **Drop Extra Columns:** Any columns in the test data but not required by the model are removed.
3. **Order Matching:** Ensures that the columns are in the same order as the model's training data.

This function ensures smooth prediction and prevents runtime errors due to feature mismatch.

```

1 def align_features(X_test, model):
2     """
3     Aligns test data features with the features expected by the model.
4
5     Parameters:
6         X_test (pd.DataFrame): Test data features.
7         model: Trained model (XGBoost or similar).
8
9     Returns:
10        pd.DataFrame: Aligned test data.
11    """
12    # Get the feature names from the model
13    model_features = model.get_booster().feature_names
14
15    # Add missing columns to X_test
16    for col in model_features:
17        if col not in X_test.columns:
18            X_test[col] = 0 # Default value for missing features
19
20    # Drop extra columns not in the model
21    X_test = X_test[model_features]
22
23    return X_test
24
25
26 def visualize_time_series(df_metadata, bldg_id, model):
27     """
28     Load and visualize time series data for energy consumption and model predictions.
29
30     Parameters:
31         df_metadata (pd.DataFrame): DataFrame containing building metadata.
32         bldg_id (str): Building ID for which the data is visualized.
33         model: Trained XGBoost model for predictions.
34     """
35    # Load the specific building data
36    file_path = f"{BUILDDINGS}/{bldg_id}.csv"
37    df_bldg = pd.read_csv(file_path)
38
39    # Clean column names

```

```

40 df_bldg.columns = [re.sub(r"^[A-Za-z0-9_]+", "_", col) for col in df_bldg.columns]
41 print(df_bldg.columns)
42
43 # Convert timestamp to datetime
44 df_bldg['timestamp'] = pd.to_datetime(df_bldg['timestamp'])
45
46 # Filter rows where minute == 0
47 df_bldg = df_bldg[df_bldg['minute'] == 0]
48
49 # Prepare features for prediction
50 X_test = df_bldg.drop(columns=['out_electricity_total_energy_consumption', 'timestamp', 'bldg_id'])
51 X_test = align_features(X_test, model)
52
53 # Predict using the model
54 df_bldg['Predicted_Energy_Consumption'] = model.predict(X_test)
55
56 # Melt the DataFrame for easier plotting
57 df_long = df_bldg.melt(
58     id_vars='timestamp',
59     value_vars=[
60         'out_electricity_total_energy_consumption',
61         'Predicted_Energy_Consumption'
62     ],
63     var_name='Measurement',
64     value_name='Value'
65 )
66
67 # Replace specific measurement names for better legend labels
68 df_long['Measurement'] = df_long['Measurement'].replace({
69     'out_electricity_total_energy_consumption': 'Actual Energy Consumption',
70     'Predicted_Energy_Consumption': 'Predicted Energy Consumption'
71 })
72
73 # Create the time series plot
74 fig = px.line(
75     df_long,
76     x='timestamp',
77     y='Value',
78     color='Measurement',
79     labels={'Value': 'Measurement Value', 'timestamp': 'Time'},
80     title=f"Time Series Data for Building ID: {bldg_id} (Actual vs Predicted)"
81 )
82
83 # Show the plot
84 fig.show()

```

✓ Visualize Time Series Data (Actual vs Predicted)

The `visualize_time_series` function provides a comprehensive view of the model's performance by comparing actual energy consumption against the predicted values.

Key components:

1. Data Preparation:

- Load the building-specific dataset and filter rows where `minute == 0` for consistent granularity.
- Align features with the trained model using the `align_features` function.

2. Predictions:

- The model predicts energy consumption using preprocessed test data.
- Predicted values are added as a new column to the dataset.

3. Visualization:


- A line plot compares actual and predicted energy consumption over time.
- Interactive legends allow users to focus on specific measurements.

This visualization helps identify patterns, trends, and areas where the model's predictions deviate from the actual values.

```
1 visualize_time_series(df_metadata, '32', model)
```

```
Index(['timestamp', 'out_electricity_total_energy_consumption',  
      'Dry_Bulb_Temperature_C_', 'Relative_Humidity_', 'heat_index', 'minute',  
      'hour', 'day', 'month', 'is_weekday', 'is_holiday', 'max_load_hourly',  
      'min_load_hourly', 'max_temp_hourly', 'min_temp_hourly', 'bldg_id'],  
      dtype='object')
```


Time Series Data for Building ID: 32 (Actual vs Predicted)



```
1 visualize_time_series(df_metadata, '105885', model)
```

```
Index(['timestamp', 'out_electricity_total_energy_consumption',  
      'Dry_Bulb_Temperature_C_', 'Relative_Humidity_', 'heat_index', 'minute',  
      'hour', 'day', 'month', 'is_weekday', 'is_holiday', 'max_load_hourly',  
      'min_load_hourly', 'max_temp_hourly', 'min_temp_hourly', 'bldg_id'],  
      dtype='object')
```

Time Series Data for Building ID: 105885 (Actual vs Predicted)




```
1 visualize_time_series(df_metadata, '1025', model)
```

```

Index(['timestamp', 'out_electricity_total_energy_consumption',
      'Dry_Bulb_Temperature_C_', 'Relative_Humidity_', 'heat_index', 'minute',
      'hour', 'day', 'month', 'is_weekday', 'is_holiday', 'max_load_hourly',
      'min_load_hourly', 'max_temp_hourly', 'min_temp_hourly', 'bldg_id'],
      dtype='object')

```

Time Series Data for Building ID: 1025 (Actual vs Predicted)



```

1 feature_importances = model.feature_importances_
2 features = []
3 features = model.get_booster().feature_names
4
5 # Create a DataFrame for better handling of the data
6 data = pd.DataFrame({
7     'Feature': features,
8     'Importance': feature_importances
9 })
10
11 # Sort features by importance for better visualization
12 data = data.sort_values(by='Importance', ascending=True)
13
14 # Plot the feature importances using Plotly
15 fig = px.bar(data, x='Importance', y='Feature', orientation='h',
16              title='Feature Importances for Load Forecasting',
17              labels={'Importance': 'Feature Importance', 'Feature': 'Feature'})
18              text='Importance')
19
20 # Improve layout for better readability
21 fig.update_layout(
22     xaxis_title="Feature Importance",
23     yaxis_title="Feature",
24     title_x=0.5, # Center the title
25     font=dict(size=12),
26     showlegend=False,
27     margin=dict(l=150, r=20, t=50, b=50), # Adjust left margin for long label
28     height=400 + 20 * len(features) # Dynamically adjust height for label siz
29 )
30
31 # Add better formatting for the text

```

```
32 fig.update_traces(texttemplate='%{text:.2f}', textposition='outside')
33
34 # Display the plot
35 fig.show()
```

